


Preparing for Assignment 4

15-418/618: Parallel Computer Architecture and Programming
Recitation. March 24, 2017


Randy Bryant

Overall Structure of System



Overall System Structure

- Clients are remote, requesting services
- Master is single-threaded
- Workers are multi-threaded. They carry out the requests
- All communication is via asynchronous messages



Master Operation

- Continually responds to *events*


- Event types

- Request from client
- Response by worker
- New worker available
- Period “tick”

- Master actions


- Send request to worker
- Send response to client
- Ask for new worker
- Kill existing worker

- Supplied code implements event manager. Your job is to handle the different event types



Worker Operation

- Maintains queue of requests from master
- Maintains pool of threads
 - Each thread runs loop
 - Get request from queue
 - Execute
 - Send response to master
- You must implement most of this



Design Features

■ Client concerns

- Different request types have different resource requirements
- Want to provide minimum latency, especially for small jobs
- Your master must control mapping of requests to workers via its scheduler

■ Resource concerns

- Demand varies over time
- Can add more workers to handle increased demand
 - Takes a while for them to come online
 - Penalized for too much resource use
- Your master must add and kill workers dynamically

Request Types

■ 418wisdom

- Compute intensive
- Low memory footprint
- Same each time

■ countprimes

- Compute intensive
- Low memory footprint
- Varies each time

■ compareprimes

- Implemented as 4 countprimes requests

■ tellmenow

- Minimal computation or memory
- Must have quick response

■ projectideas

- Fills entire L3 cache

Benchmarking Study of Request Types

■ All running on GHC machine

- 8 cores, each 2x hyperthreaded
- 20MB L3 cache
- Latedays will be slightly different

■ Methodology

- Code in

`/afs/cs.cmu.edu/academic/class/15418-s17/recitations/recw8/recw8-code/asst4_bench`

- Request some combination of different jobs
- Possibly restrict number of outstanding jobs of given type

■ Measurements

- For each job type: Minimum, Average, Maximum elapsed time
- Overall elapsed time
- All times in milliseconds

Benchmark Run Example

16 threads				
	wisdom	primes	tellme	project
Jobs:	64	0	0	4
Limits:	64	0	0	1
Min:	1056	--	--	3978
avg:	1229	--	--	4335
Max:	1343	--	--	5161
Total time: 17343ms				

- 16 threads
- Jobs: 64 418wisdom + 4 projectideas
- Limits: None for 418wisdom, 1 at a time for projectideas
- Show min/avg/max for each type
- Overall time

418wisdom

- 64 jobs
- Vary number of threads
- No limit on number of threads running requests simultaneously

Threads	Average	Total
8	1080	8687
16	1243	5037
32	2274	5142
64	4342	5041

- Questions (remember 8 cores, each 2x hyperthreaded)
 - Why does average increase with number of threads?
 - Why does overall time improve going from 8 to 16 threads?
 - Why does it stay flat beyond 16 threads?

Compute Intensive Requests

- 64 jobs
- 16 threads

	418wisdom	countprimes	tellmenow
Minimum	1207	2	< 1
Average	1245	1254	< 1
Maximum	1344	2718	< 1

- **Questions:**
 - What do these numbers imply about scheduling?

projectideas

- 4 jobs
- 16 threads
- Limit number of simultaneous jobs

Limit	Average	Total
1	4028	16116
2	10473	21077
4	14293	15027

- **Questions:**
 - Why does both the average and the total increase when going from 1 to 2 running simultaneously?
 - Why does the total drop when run 4 simultaneously?
 - What would be a good scheduling policy?

Mixing 418wisdom and projectideas

- Combine 418wisdom jobs + projectideas jobs
- 16 threads
- Limit of 1 simultaneous projectideas job

projectideas jobs	418wisdom jobs	418wisdom average	projectideas average	Total
0	64	1243	—	5037
4	0	—	4028	16116
4	64	1229	4151	16951

■ Questions:

- Why is the average time for 418wisdom unaffected by projectideas?
- Why is the average time for projectideas unaffected by 418wisdom?
- What does this imply about scheduling policies?

Useful Features of C++

- **Reference parameters**
- **C++ strings**
- **Class basics**
- **Templates**
- **Memory management**

[/afs/cs.cmu.edu/academic/class/15418-s17/recitations/recw8/recw8-code/cpp](http://afs/cs.cmu.edu/academic/class/15418-s17/recitations/recw8/recw8-code/cpp)

Reference Parameters (eg1-refparam)

```
// Demonstration of reference parameters
#include <iostream>

// C style
int pincr(int *xp) {
    int r = *xp;
    (*xp)++;
    return r;
}

// C++ style
int rincr(int &x) {
    int r = x;
    x++;
    return r;
}

int main(int argc, char *argv[]) {
    int x = 1;
    int y = pincr(&x);
    int z = rincr(x);
    std::cout << "x=" << x << ", y=" << y
                << ", z=" << z << std::endl;
    return 0;
}
```

- Function rincr uses call-by-reference
- What are resulting values of x, y, and z?

Implementation of Reference Parameters

C Pointer Code

```
// C style
int pincr(int *xp) {
    int r = *xp;
    (*xp)++;
    return r;
}
```

C++ Reference Code

```
// C++ style
int rincr(int &x) {
    int r = x;
    x++;
    return r;
}
```

```
_Z5pincrPi:
    movl    (%rdi), %eax
    leal    1(%rax), %edx
    movl    %edx, (%rdi)
    ret
```

```
_Z5rincrRi:
    movl    (%rdi), %eax
    leal    1(%rax), %edx
    movl    %edx, (%rdi)
    ret
```


C++ strings (eg2-string)

```
// Demonstration of C++ strings
#include <string>
#include <iostream>
#include <stdio.h>

std::string stringify(char *s) {
    std::string ss = s;
    return ss;
}

int main(int arg, char *argv[]) {
    char buf[10];
    std::string strings[5];
    for (int i = 0; i < 5; i++) {
        sprintf(buf, "i=%d", i);
        strings[i] = stringify(buf);
    }
    for (int i = 0; i < 5; i++) {
        std::cout << "String " << i << ":"
                  << strings[i] << std::endl;
    }
    return 0;
}
```

- C++ provides managed strings
- Automatic garbage collection via reference counting
- How is this possible, given rest of language requires explicit memory management?

Comparing to C strings (eg2a-string)

```
// Comparison with C strings
// Does this code work?
#include <iostream>
#include <stdio.h>
#include <string.h>

char *stringify(char *s) {
    char buf[10];
    strcpy(buf, s);
    return buf;
}

int main(int arg, char *argv[]) {
    char buf[10];
    char *strings[5];
    for (int i = 0; i < 5; i++) {
        sprintf(buf, "i=%d", i);
        strings[i] = stringify(buf);
    }
    for (int i = 0; i < 5; i++) {
        std::cout << "String " << i << ":"
        << strings[i] << std::endl;
    }
    return 0;
}
```

- C strings are simply arrays of characters
- No management
- Programmer must explicitly manage storage

Implementation of string-returning function

```
// Nominal version
std::string stringify(char *s) {
    std::string ss = s;
    return ss;
}

// True implementation
void rstringify(char *s, std::string &ss) {
    ss = s;
}
```

Returning struct in C

- Caller allocates space on stack
- Passes pointer to the callee
- Callee fills in fields

Returning object in C++

- Same strategy

Class Example (eg3-class)

```
// Implementation of linked list of integers
#include <iostream>

class ListEle {
private:
    int val;
    ListEle *next;

public:
    ListEle(int v, ListEle *np)
        { val = v; next = np; }

    ListEle *getNext() { return next; }

    int getValue() { return val; }
};
```

- C++ class is a variant of a struct
 - Reference with “.” or “->” operator
- Can make some fields inaccessible outside of class code
- Can write one or more constructor functions

Class Example (cont)

```
class List {  
private:  
    ListEle *head;  
  
public:  
    List() { head = NULL; }  
  
    void insert(int v) { head = new ListEle(v, head); }  
  
    int front() {  
        if (head) return head->getValue();  
        else      return -1;  
    }  
  
    void pop() {  
        if (head) head = head->getNext();  
    }  
  
    bool isEmpty() { return head == NULL; }  
};
```

- Allocate element with *new* operation
- Can you find the memory leak?
- How would you fix it? (see eg3a-class.cpp)

Using Class

```
int main(int arg, char *argv[]) {
    List ls;
    for (int i = 0; i < 5; i++) {
        ls.insert(i);
    }

    while (!ls.isEmpty()) {
        int v = ls.front();
        std::cout << "Popped value " << v << std::endl;
        ls.pop();
    }
    return 0;
}
```

- List head allocated on stack
- List elements allocated on heap

Template Example (eg4-template)

```
// Implementation of linked list of integers
#include <iostream>

template <class T>
class ListEle {
private:
    T val;
    ListEle *next;
public:

    ListEle(T v, ListEle *np)
    { val = v; next = np; }

    ListEle *getNext() { return next; }

    T getValue() { return val; }
};
```

- Template allows generic code that is later instantiated for one or more types

Templates (eg4-template)

```
template <class T>
class List {
private:
    ListEle<T> *head;
public:
    List() { head = NULL; }

    void insert(T v) { head = new ListEle<T>(v, head); }

    T front() { return head->getValue(); }

    void pop() {
        if (head) {
            ListEle<T> *save = head;
            head = head->getNext();
            delete save;
        }
    }

    bool isEmpty() { return head == NULL; }
};
```

- Use T in places where previously had int
- Note use of *delete* to avoid memory leak

Template instantiation (eg4-template)

```
int main(int arg, char *argv[]) {
    List<int> ls;
    for (int i = 0; i < 5; i++) {
        ls.insert(i);
    }

    while (!ls.isEmpty()) {
        int v = ls.front();
        std::cout << "Popped value " << v << std::endl;
        ls.pop();
    }
    return 0;
}
```

- Creates list of integers
- But, can use for list of any object type

Freeing storage with destructors (eg5-deallocate)

Destructor for ListEle

```
~ListEle() {  
    std::cout << "Destroying element with value " << val << std::endl;  
    if (next)  
        delete next;  
}
```

Destructor for List

```
~List() { if (head) delete head; }
```

- Called when object deallocated
- Stack-allocated objects deallocated when exit scope of declaration
- Heap-allocated objects only deallocated by *delete*

Unintended deallocation (eg5a-deallocate)

```
int main(int arg, char *argv[]) {
    List ls;
    for (int i = 0; i < 5; i++) {
        ls.insert(i);
    }

    while (!ls.isEmpty()) {
        int v = ls.front();
        std::cout << "Popped value " << v << std::endl;
        ls.pop();
    }
    return 0;
}
```

- What happens when call `ls.pop()`?
- Why doesn't this happen in garbage-collected languages?

```
void pop() {
    if (head) {
        ListEle *save = head;
        head = head->getNext();
        delete save;
    }
}
```

Controlling deallocation (eg5b-deallocate)

```
void pop() {  
    if (head) {  
        ListEle *save = head;  
        head = head->getNext();  
        save->unlink();  
        delete save;  
    }  
}
```

- Set next pointer to null before deleting

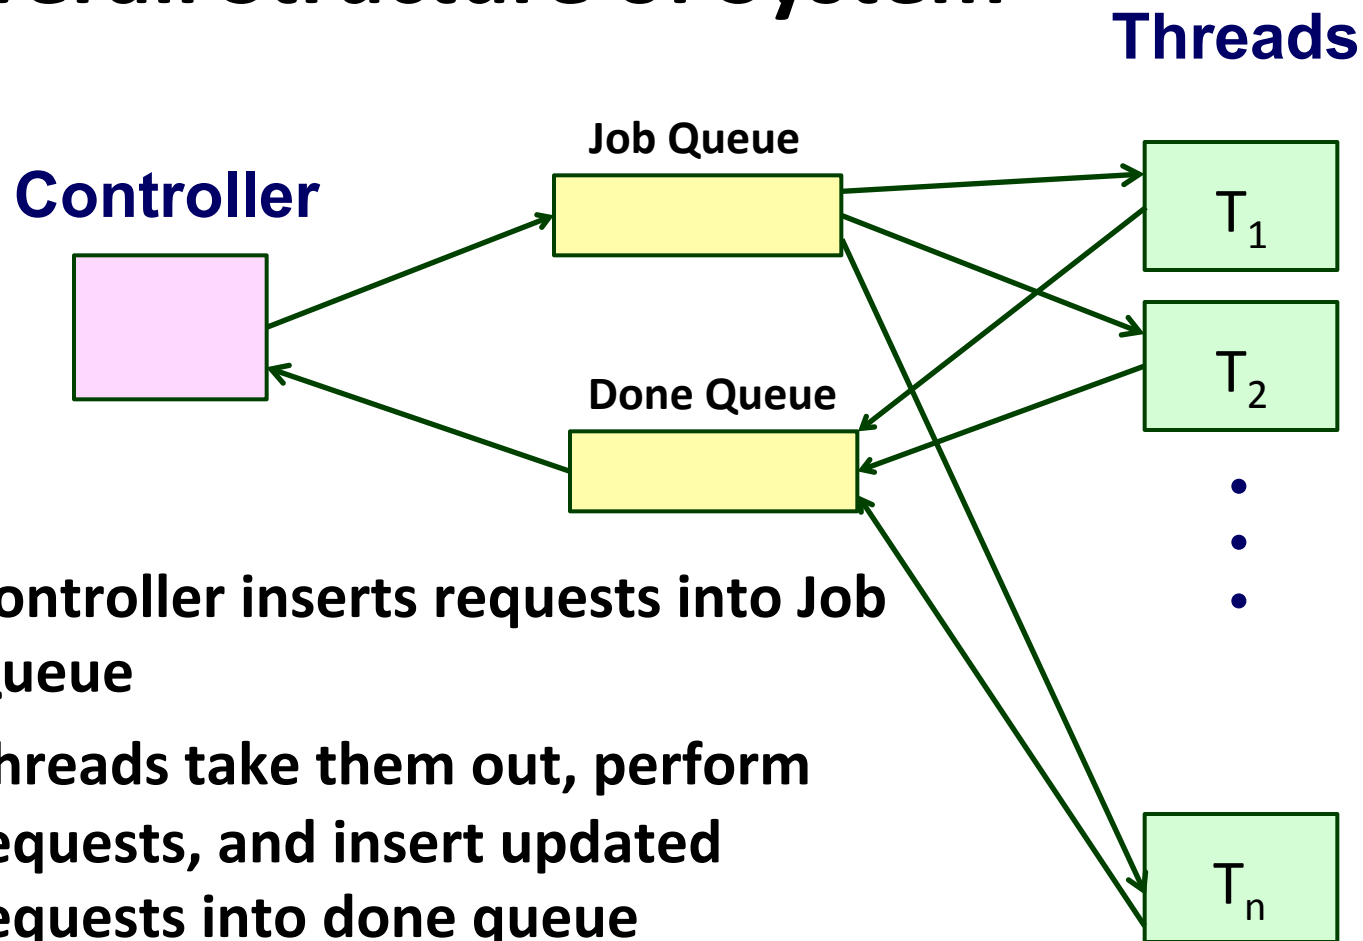
Program Example

- **The benchmarking framework described earlier**

[/afs/cs.cmu.edu/academic/class/15418-s17/recitations/recw8/recw8-code/asst4_bench](https://afs.cs.cmu.edu/academic/class/15418-s17/recitations/recw8/recw8-code/asst4_bench)

- **Elements similar to those required for assignment worker**
 - But, much simpler

Overall Structure of System



- Controller inserts requests into Job Queue
- Threads take them out, perform requests, and insert updated requests into done queue
- Controller reads completed requests
- All synchronization through queues

Request format (run.cpp)

```
// Record to track single job.
// Continues from creation to execution to completion
class Job {
    static int nextjid; // Use for generating unique ids

public:
    int benchmark; // Which benchmark type
    int id;        // Unique ID for job
    int msec;      // How many milliseconds did it require

    Job(int b) {
        id = nextjid++;
        benchmark = b;
        msec = 0;
    }
};

int Job::nextjid; // Initializes to 0
```

- Integer benchmark indicates job type
- msec will be filled in by worker
- Variable nextjid part of class, not object
 - Use to generate unique IDs

Queues

```
template <class T>
class WorkQueue {
private:
    ...
public:

    WorkQueue() {
        ...
    }

    T get_work() {
        ...
    }

    void put_work(const T& item) {
        ...
    };
};
```

Code Samples

```
// Maintain two job queues
typedef WorkQueue<Job> JobQueue;

// Jobs waiting to be executed
JobQueue *newJobQueue = NULL;
// Jobs that have completed
JobQueue *doneQueue = NULL;
```

```
newJobQueue = new JobQueue;
doneQueue = new JobQueue;
```

```
Job job = Job(b);
newJobQueue->put_work(job);
```


```
Job job = doneQueue->get_work();
int b = job.benchmark;
report_job(b, job.msecs);
```

- Make use of code provided in `work_queue.h`
- Uses mutex and condition variable to ensure safe insert/remove

Termination


Threads

Controller



Job Queue

(Empty)





Done Queue




Thread code

```
Job job = newJobQueue->get_work();
```



⋮



- Controller can detect when done by counting jobs removed from done queue
- But, threads hang up trying to remove from job queue
- How can they be notified?

Logging

- Google logging code included as part of assignment code

```
DLOG(INFO) << "Thread #" << tid << " working on request [" << request_tag  
    << ":" << request_string << "]\n";
```

- Logging information accumulated in separate log file for each run
- Tip: Make up useful prefixes to allow extraction of interesting lines with grep

```
DLOG(INFO) << "TCK\tTIME " << mstate.num_seconds  
    << ". Queued requests:\n";
```

Some Advice

■ Lots of design thinking required

- Data structures used by master
- How to manage jobs with different priorities
- Implementing elasticity
- Priority scheduler for worker
 - May need to extend provided queue code
 - But, it's tricky stuff

■ Expectations (Lines of code)

- master: 132 (provided) → 601 (My solution)
- worker: 94 (provided) → 235 (My solution)
- Provided code is mostly comments

■ Harsh reality

- The latedays machines will be horrendously overloaded as approach deadline